



# **CS 412 Intro. to Data Mining**

**Chapter 7 : Advanced Frequent Pattern Mining**

**Qi Li, Computer Science, Univ. Illinois at Urbana-Champaign, 2018**



# **Chapter 7 : Advanced Frequent Pattern Mining**

---

- Mining Diverse Patterns
- Constraint-Based Frequent Pattern Mining
- Sequential Pattern Mining
- Graph Pattern Mining
- Pattern Mining Application: Mining Software Copy-and-Paste Bugs
- Summary



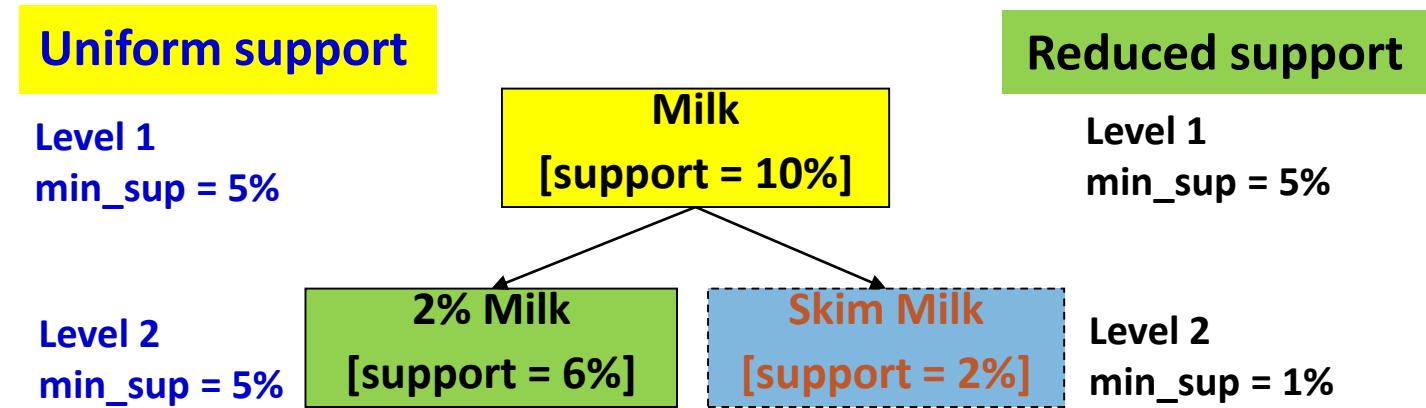
# Mining Diverse Patterns

---

- Mining Multiple-Level Associations
- Mining Multi-Dimensional Associations
- Mining Quantitative Associations
- Mining Negative Correlations
- Mining Compressed and Redundancy-Aware Patterns

# Mining Multiple-Level Frequent Patterns

- ❑ Min-support thresholds for hierarchy items
- ❑ **Uniform** min-support across multiple levels (reasonable?)
- ❑ **Level-reduced** min-support: Items at the lower level are expected to have lower support
- ❑ Efficient mining: *Shared* multi-level mining
- ❑ Use the lowest min-support to pass down the set of candidates



# Redundancy Filtering at Mining Multi-Level Associations

---

- Redundancy filtering: redundant due to “ancestor” relationships
  - milk  $\Rightarrow$  wheat bread [support = 8%, confidence = 70%] (1)
  - 2% milk  $\Rightarrow$  wheat bread [support = 2%, confidence = 72%] (2)
  - Suppose the 2% milk sold is about  $\frac{1}{4}$  of milk sold in gallons
    - (2) should be able to be “derived” from (1)

# Redundancy Filtering at Mining Multi-Level Associations

---

- milk  $\Rightarrow$  wheat bread [support = 8%, confidence = 70%] (1)
- 2% milk  $\Rightarrow$  wheat bread [support = 2%, confidence = 72%] (2)
- A rule is *redundant* if its support is close to the “expected” value, according to its “ancestor” rule, and it has a similar confidence as its “ancestor”
- Rule (1) is an ancestor of rule (2), which one to prune?

# Customized Min-Supports for Different Kinds of Items

---

- ❑ Same min-support threshold **for all** so far
- ❑ Diamonds, watches: valuable but **less** frequent
- ❑ One Method: Use **group-based** “individualized” min-support
  - ❑ E.g., {diamond, watch}: 0.05%; {bread, milk}: 5%; ...
  - ❑ How to mine such rules efficiently?
  - ❑ Existing scalable mining algorithms can be easily extended to cover such cases

# Mining Multi-Dimensional Associations

---

- Single-dimensional rules (e.g., items are all in “product” dimension)
  - $\text{buys}(X, \text{“milk”}) \Rightarrow \text{buys}(X, \text{“bread”})$
- Multi-dimensional rules (i.e., items in  $\geq 2$  dimensions or predicates)
  - Inter-dimension association rules (*no repeated predicates*)
    - $\text{age}(X, \text{“18-25”}) \wedge \text{occupation}(X, \text{“student”}) \Rightarrow \text{buys}(X, \text{“coke”})$
  - Hybrid-dimension association rules (*repeated predicates*)
    - $\text{age}(X, \text{“18-25”}) \wedge \text{buys}(X, \text{“popcorn”}) \Rightarrow \text{buys}(X, \text{“coke”})$
- Attributes can be categorical or numerical
  - Categorical Attributes (e.g., *profession*, *product*: no ordering among values): Data cube for inter-dimension association
  - Quantitative Attributes: Numeric, implicit ordering among values—discretization, clustering, and gradient approaches

# Mining Quantitative Associations

---

- Mining associations with numerical attributes
  - E.g.: Numerical attributes: **age** and **salary**
- Methods
  - **Static discretization** based on predefined concept hierarchies
    - Discretization on each dimension with hierarchy
      - age: {0-10, 10-20, ..., 90-100} → {young, mid-aged, old}
  - **Dynamic discretization** based on data distribution
  - **Clustering**: Distance-based association
    - First one-dimensional clustering, then association
  - **Deviation analysis**:
    - Gender = female ⇒ Wage: mean=\$7/hr (overall mean = \$9)

# Mining Extraordinary Phenomena in Quantitative Association Mining

---

- Mining extraordinary (i.e., interesting) phenomena
  - E.g.: **Gender = female**  $\Rightarrow$  **Wage**: mean=\$7/hr (overall mean = \$9)
  - **LHS**: a subset of the population
  - **RHS**: an extraordinary behavior of this subset
- The rule is accepted only if a statistical test (e.g., Z-test) confirms the inference with high confidence
  
- Subrule: Highlights the extraordinary behavior of a subset of the population of the super rule
  - E.g.: **(Gender = female) ^ (South = yes)**  $\Rightarrow$  mean wage = \$6.3/hr
- Rule condition can be categorical or numerical (quantitative rules)
  - E.g.: **Education in [14-18] (yrs)**  $\Rightarrow$  mean wage = \$11.64/hr

# Rare Patterns

---

- Rare patterns
  - Very low support but interesting (e.g., buying Rolex watches)
  - How to mine them? Setting individualized, group-based min-support thresholds for different groups of items

# Negative Patterns

---

- ❑ Negative patterns
  - ❑ Negatively correlated: Unlikely to happen together
  - ❑ Ex.: Since it is unlikely that the same customer buys both a **Ford Expedition** (an SUV car) and a **Ford Fusion** (a hybrid car), buying a **Ford Expedition** and buying a **Ford Fusion** are likely negatively correlated patterns
  - ❑ How to define negative patterns?
- ❑ A support-based definition of negative correlated patterns
  - ❑ If itemsets A and B are **both frequent** but rarely occur together, i.e.,  $\text{sup}(A \cup B) \ll \text{sup}(A) \times \text{sup}(B)$

Does this remind you the definition of *lift*?

# Defining Negative Correlated Patterns

---

Is this a good definition for large transaction datasets?

- Ex.: Suppose a store sold two needle packages A and B 100 times each, but only one transaction contained both A and B
  - When there are in total 200 transactions, we have
    - $s(A \cup B) = 0.005, s(A) \times s(B) = 0.25, s(A \cup B) << s(A) \times s(B)$
  - But when there are  $10^5$  transactions, we have
    - $s(A \cup B) = 1/10^5, s(A) \times s(B) = 1/10^3 \times 1/10^3, s(A \cup B) > s(A) \times s(B)$
- What is the problem?—Null transactions: The support-based definition is not null-invariant!

# Defining Negative Correlation: Need Null-Invariance in Definition

- A Kulczynski measure-based definition
  - If itemsets A and B are frequent but  $(s(A \cup B)/s(A) + s(A \cup B)/s(B))/2 < \epsilon$ ,  
then A and B are negatively correlated
- For the same needle package problem:
  - No matter there are in total 200 or  $10^5$  transactions
  - If  $\epsilon = 0.01$ , we have  
$$(s(A \cup B)/s(A) + s(A \cup B)/s(B))/2 = (0.01 + 0.01)/2 < \epsilon$$

negative pattern threshold

# Mining Compressed Patterns

Pat-ID	Item-Sets	Support
P1	{38,16,18,12}	205227
P2	{38,16,18,12,17}	205211
P3	{39,38,16,18,12,17}	101758
P4	{39,16,18,12,17}	161563
P5	{39,16,18,12}	161576

- Closed patterns
  - P1, P2, P3, P4, P5
  - Emphasizes too much on support
- Max-patterns
  - P3: information loss
- Desired output (a good balance):
  - P2, P3, P4

- Why mining compressed patterns? Too many scattered patterns but not so meaningful

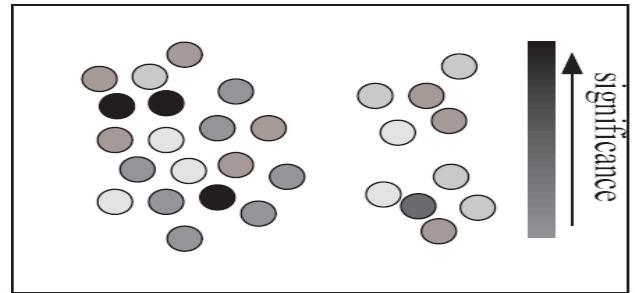
- Pattern distance measure

$$Dist(P_1, P_2) = 1 - \frac{|T(P_1) \cap T(P_2)|}{|T(P_1) \cup T(P_2)|}$$

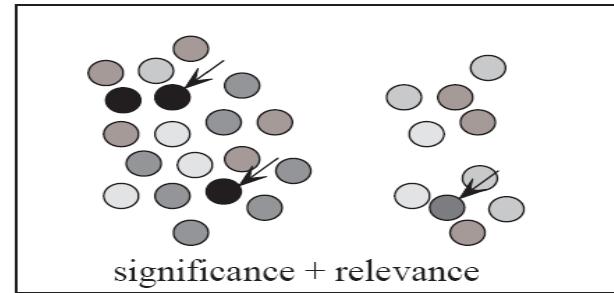
- $\delta$ -clustering: For each pattern P, find all patterns which can be expressed by P and whose distance to P is within  $\delta$  ( $\delta$ -cover)
  - All patterns in the cluster can be represented by P

# Redundancy-Aware Top-k Patterns

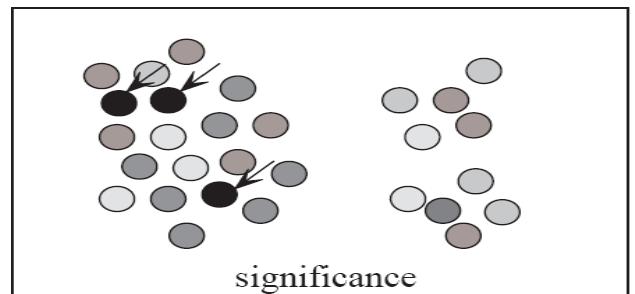
- Desired patterns: high significance & low redundancy



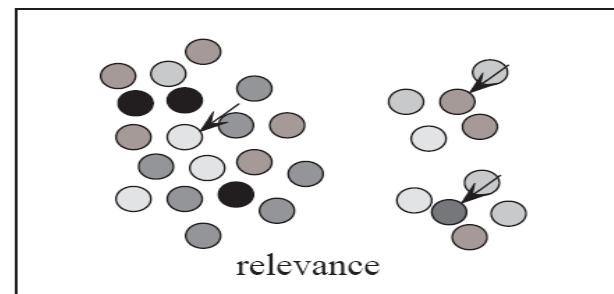
(a) a set of patterns



(b) redundancy-aware top- $k$



(c) traditional top- $k$



(d) summarization

- Method: Use MMS (Maximal Marginal Significance) for measuring the combined significance of a pattern set
- Xin et al., Extracting Redundancy-Aware Top-K Patterns, KDD'06

# **Chapter 7 : Advanced Frequent Pattern Mining**

---

- ❑ Mining Diverse Patterns
- ❑ Constraint-Based Frequent Pattern Mining
- ❑ Sequential Pattern Mining
- ❑ Graph Pattern Mining
- ❑ Pattern Mining Application: Mining Software Copy-and-Paste Bugs
- ❑ Summary



# Constraint-Based Pattern Mining

---

- Why Constraint-Based Mining?
- Different Kinds of Constraints: Different Pruning Strategies
- Constrained Mining with Pattern Anti-Monotonicity
- Constrained Mining with Pattern Monotonicity
- Constrained Mining with Convertible Constraints
- Constrained Mining with Data Anti-Monotonicity
- Constrained Mining with Succinct Constraints
- Handling Multiple Constraints

# Why Constraint-Based Mining?

---

- Pattern mining in practice: Often a user-guided, **interactive** process
  - User directs what to be mined using a **data mining query language** (or a graphical user interface), **specifying various kinds of constraints**
- What is constraint-based mining?
  - Mine together with user-provided constraints
- Why constraint-based mining?
  - User flexibility: User provides **constraints** on what to be mined
  - Optimization: System explores such constraints for mining efficiency
  - E.g., Push constraints deeply into the mining process

# Various Kinds of User-Specified Constraints in Data Mining

---

- **Knowledge type constraint**—Specifying what kinds of knowledge to mine
  - E.g.: Classification, association, clustering, outlier finding, ...
- **Data constraint**—using SQL-like queries
  - E.g.: Find products sold together in **NY** stores **this year**
- **Dimension/level constraint**—similar to projection in relational database
  - E.g.: In relevance to **region**, **price**, **brand**, **customer category**
- **Interestingness constraint**—various kinds of thresholds
  - E.g.: Strong rules:  $\text{min\_sup} \geq 0.02$ ,  $\text{min\_conf} \geq 0.6$ ,  $\text{min\_correlation} \geq 0.7$
- **Rule (or pattern) constraint**  **The focus of this study**
  - E.g.: Small sales (price < \$10) triggers big sales (sum > \$200)

# Pattern Space Pruning with Pattern Anti-Monotonicity

Item	Price	Profit
a	100	40
b	40	0
c	150	-20
d	35	-15
e	55	-30
f	45	-10
g	80	20
h	10	5

Note: item.price > 0  
Profit can be negative

- A constraint  $c$  is **anti-monotone**
  - If an itemset  $S$  **violates** constraint  $c$ , so does any of its superset
  - That is, mining on itemset  $S$  can be terminated
- E.g. 1:  $c_1: \text{sum}(S.\text{price}) \leq v$  is **anti-monotone**
  - Sum grows as you add more items
- E.g. 2:  $c_2: \text{range}(S.\text{profit}) \leq 15$  is **anti-monotone**
  - Itemset  $ab$  violates  $c_2$  ( $\text{range}(ab) = 40$ )
  - So does every superset of  $ab$

# Pattern Space Pruning with Pattern Anti-Monotonicity

TID	Transaction
10	a, b, c, d, f, h
20	b, c, d, f, g, h
30	b, c, d, f, g
40	a, c, e, f, g

min_sup = 2		
Item	Price	Profit
a	100	40
b	40	0
c	150	-20
d	35	-15
e	55	-30
f	45	-10
g	80	20
h	10	5

- E.g. 3.  $c_3: \text{sum}(S.\text{Price}) \geq v$  is not anti-monotone
- E.g. 4. Is  $c_4: \text{support}(S) \geq \sigma$  anti-monotone?
  - Yes! Apriori pruning is essentially pruning with an anti-monotonic constraint!

# Pattern Monotonicity and Its Roles

TID	Transaction
10	a, b, c, d, f, h
20	b, c, d, f, g, h
30	b, c, d, f, g
40	a, c, e, f, g

min\_sup = 2

Item	Price	Profit
a	100	40
b	40	0
c	150	-20
d	35	-15
e	55	-30
f	45	-10
g	80	20
h	10	5

- A constraint  $c$  is *monotone*: If an itemset  $S$  **satisfies** the constraint  $c$ , so does any of its superset
  - That is, we do not need to check  $c$  in subsequent mining
  - Not as beneficial as anti-monotone
- E.g. 1:  $c_1: \text{sum}(S.\text{Price}) \geq v$  is **monotone**
- E.g. 2:  $c_2: \text{min}(S.\text{Price}) \leq v$  is **monotone**
- E.g. 3:  $c_3: \text{range}(S.\text{profit}) \geq 15$  is **monotone**
  - Itemset  $ab$  satisfies  $c_3$
  - So does every superset of  $ab$

# Apriori for Pattern Anti-Monotone Constraint

Item	Price
1	1
2	2
3	3
4	4
5	5

Database D

TID	Items
10	1 3 4
20	2 3 5
30	1 2 3 5
40	2 5

Scan D

itemset	sup.
{1}	2
{2}	3
{3}	3
{4}	1
{5}	3

$F_1$

itemset	sup.
{1}	2
{2}	3
{3}	3
{5}	3

Can be  
chopped  
early



$C_2$

itemset	sup.
{1 2}	1
{1 3}	2
{1 5}	1
{2 3}	2
{2 5}	3
{3 5}	2

$C_2$

itemset
{1 2}
{1 3}
{1 5}
{2 3}
{2 5}
{3 5}

Scan D

$F_2$

itemset	sup.
{1 3}	2
{2 3}	2
{2 5}	3
{3 5}	2



itemset
{2 3 5}

Scan D

itemset	sup.
{2 3 5}	2

Min\_sup=2

Constraint:

$\text{Sum}\{\text{S.price}\} < 5$

# Convertible Constraints: Ordering Data in Transactions

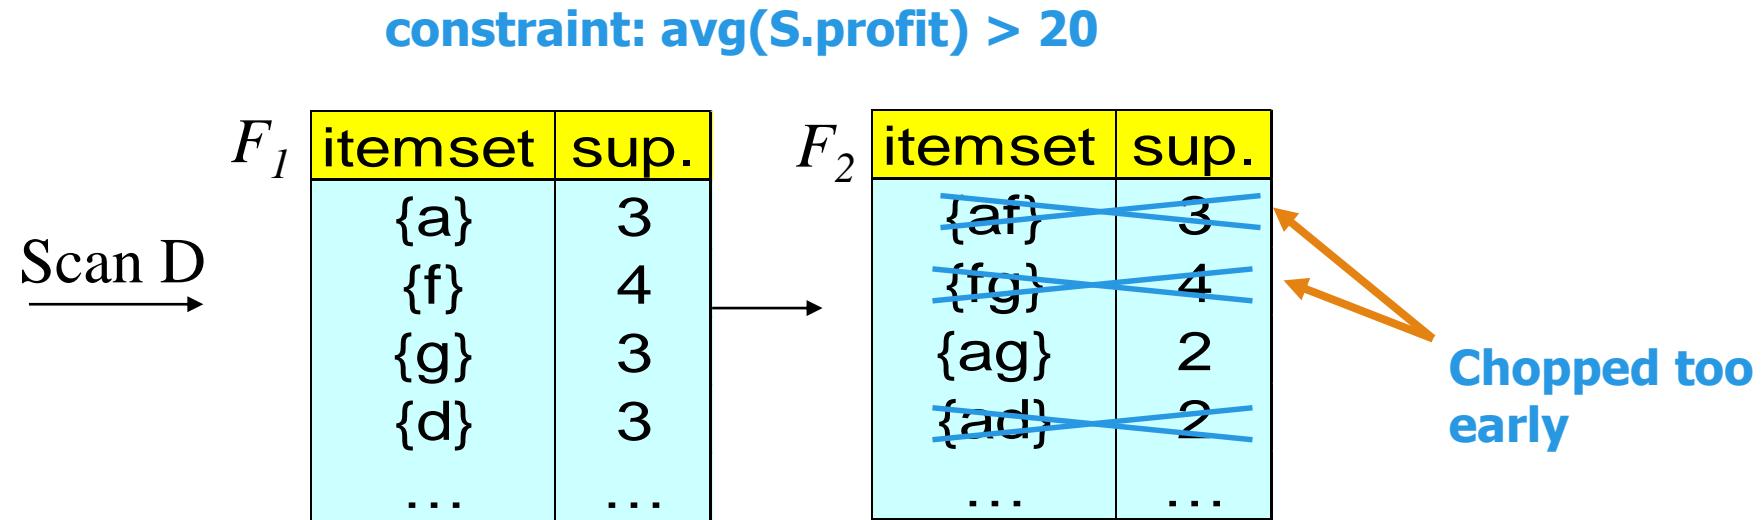
TID	Transaction
10	a, b, c, d, f, h
20	a, b, c, d, f, g, h
30	b, c, d, f, g
40	a, c, e, f, g

Item	Price	Profit
a	100	40
b	40	0
c	150	-20
d	35	-15
e	55	-30
f	45	-5
g	80	30
h	10	5

- Convert tough constraints into (anti-)monotone by proper ordering of items in transactions
- Examine  $c_1: \text{avg}(S.\text{profit}) > 20$ 
  - Order items in (profit) value-descending order
    - $\langle a, g, h, b, f, d, c, e \rangle$
    - An itemset  $ab$  violates  $c_1$  ( $\text{avg}(ab) = 20$ )
    - So does  $ab^*$  (i.e.,  $ab$ -projected DB)
    - $C_1$ : anti-monotone if patterns grow in the right order!

# Can item-reordering work for Apriori?

TID	Transaction	
10	a, b, c, d, f, h	
20	a, b, c, d, f, g, h	
30	b, c, d, f, g	
40	a, c, e, f, g	
min_sup = 2		
Item	Price	Profit
a	100	40
b	40	0
c	150	-20
d	35	-15
e	55	-30
f	45	-5
g	80	30
h	10	5



- $\text{avg}(gf) = 12.5 < 20$ ,  $\text{avg}(af) = 17.5 < 20$ ,  $\text{avg}(ag) = 35 > 20$
- But  $\text{avg}(agf) = 21.7 > 20$
- Apriori will not generate “agf” as a candidate

# Data Space Pruning with Data Anti-Monotonicity

TID	Transaction
10	a, b, c, d, f, h
20	b, c, d, f, g, h
30	b, c, d, f, g
40	a, c, e, f, g

min_sup = 2		
Item	Price	Profit
a	100	40
b	40	0
c	150	-20
d	35	-15
e	55	-30
f	45	-10
g	80	20
h	10	5

- ❑ A constraint  $c$  is *data anti-monotone*: In the mining process, if a data entry  $t$  cannot contribute to a pattern  $p$  satisfying  $c$ ,  $t$  cannot contribute to  $p$ 's superset either
- ❑ Data space pruning: Data entry  $t$  can be pruned
- ❑ E.g. 1:  $c_1: \text{sum}(S.\text{Profit}) \geq v$  is *data anti-monotone*
  - ❑ Let constraint  $c_1$  be:  $\text{sum}(S.\text{Profit}) \geq 25$
  - ❑  $T_{30}: \{b, c, d, f, g\}$  can be removed since none of their combinations can make an  $S$  whose sum of the profit is  $\geq 25$

# Data Space Pruning with Data Anti-Monotonicity

TID	Transaction
10	a, b, c, d, f, h
20	b, c, d, f, g, h
30	b, c, d, f, g
40	a, c, e, f, g

min_sup = 2		
Item	Price	Profit
a	100	40
b	40	0
c	150	-20
d	35	-15
e	55	-30
f	45	-10
g	80	20
h	10	5

- ❑ A constraint  $c$  is **data anti-monotone**: In the mining process, if a data entry  $t$  cannot contribute to a pattern  $p$  satisfying  $c$ ,  $t$  cannot contribute to  $p$ 's superset either
- ❑ Data space pruning: Data entry  $t$  can be pruned
  - ❑ E.g. 2:  $c_2: \min(S.\text{Price}) \leq v$  is **data anti-monotone**
    - ❑ Consider  $v = 5$  but every item in a transaction, say  $T_{50}$ , has a price higher than 10
  - ❑ E.g. 3:  $c_3: \text{range}(S.\text{Profit}) > 25$  is **data anti-monotone**

# Data Space Pruning Should Be Explored Recursively

TID	Transaction	Item	Profit
10	a, b, c, d, f, h	a	40
20	b, c, d, f, g, h	b	0
30	b, c, d, f, g	c	-20
40	a, c, e, f, g	d	-15
min_sup = 2		e	-30
		f	-10
		g	20
		h	5

- ❑ Example.  $c_3: \text{range}(S.\text{Profit}) > 25$
- ❑ We check b's projected database
- ❑ But item “a” is infrequent ( $\text{sup} = 1$ )
- ❑ After removing “a (40)” from  $T_{10}$
- ❑  $T_{10}$  cannot satisfy  $c_3$  any more
  - ❑ Since “b (0)”, “c (-20)”, “d (-15)”, “f (-10)”, “h (5)”
- ❑ By removing  $T_{10}$ , we can also prune “h” in  $T_{20}$

b's-proj. DB



TID	Transaction
10	a, c, d, f, h
20	c, d, f, g, h
30	c, d, f, g

b's-proj. DB

TID	Transaction
10	c, d, f, h
20	c, d, f, g, h
30	c, d, f, g

# Data Space Pruning Should Be Explored Recursively

TID	Transaction
10	a, c, d, f, h
20	c, d, f, g, h
30	c, d, f, g

Recursive  
Data  
Pruning

single branch: cdfg: 2

Constraint:  
 $\text{range}\{\text{S}.profit\} > 25$

Only a single branch “cdfg: 2”  
to be mined in b’s projected DB

- ❑ Note:  $c_3$  prunes  $T_{10}$  effectively only after “a” is pruned (by min-sup) in b’s projected DB

# Succinctness: Pruning Both Data and Pattern Spaces

---

- Succinctness: If the constraint  $c$  can be enforced by directly manipulating the data
- E.g. 1: To find those patterns containing item  $i$ 
  - Mine only  $i$ -projected DB (data space pruning)
- E.g. 2: To find those patterns without item  $i$ 
  - Remove  $i$  from DB and then mine (pattern space pruning)
- E.g. 3:  $c_3: \min(S.\text{Price}) \leq v$  is succinct
  - Start with only items whose price  $\leq v$  and remove transactions with high-price items only (pattern + data space pruning)
- E.g. 4:  $c_4: \sum(S.\text{Price}) \geq v$  is not succinct
  - It cannot be determined beforehand since sum of the price of itemset  $S$  keeps increasing

# Apriori + Succinct Constraint

Item	Price
1	1
2	2
3	3
4	4
5	5

Database D

TID	Items
10	1 3 4
20	2 3 5
30	1 2 3 5
40	2 5

Scan D

itemset	sup.
{1}	2
{2}	3
{3}	3
{4}	1
{5}	3

$F_1$

itemset	sup.
{1}	2
{2}	3
{3}	3
{4}	3
{5}	3

Min\_sup=2

Constraint:

$\min\{S.\text{price}\} \leq 2$

Chopped too early

Database D

TID	Items
10	1 3 4
20	2 3 5
30	1 2 3 5
40	2 5

Scan D

itemset	sup.
{1}	2
{2}	1
{3}	2
{4}	1
{5}	1

$F_1$

itemset	sup.
{1}	2
{3}	2

Min\_sup=2

Constraint:

$\min\{S.\text{price}\} \leq 1$

# Constrained FP-Growth: Push a Succinct Constraint Deep

Item	Price
1	1
2	2
3	3
4	4
5	5

TID	Items
10	1 3 4
20	2 3 5
30	1 2 3 5
40	2 5

Remove infrequent length 1

TID	Items
10	1 3
20	2 3 5
30	1 2 3 5
40	2 5

Min\_sup=2

Constraint:

$\min\{S.\text{price}\} \leq 2$

No Need to project on 3 or 5

1-Projected DB

TID	Items
10	3 4
30	2 3 5

2-Projected DB

TID	Items
20	3 5
30	1 3 5
40	5

# Different Kinds of Constraints Lead to Different Pruning Strategies

- In summary, constraints can be categorized as **pattern space pruning** constraints vs. **data space pruning** constraints

Pattern space pruning constraints	Data space pruning constraints
<ul style="list-style-type: none"><li><b>Anti-monotonic:</b> If constraint <math>c</math> is violated, its further mining can be terminated</li><li><b>Monotonic:</b> If <math>c</math> is satisfied, no need to check <math>c</math> again</li><li><b>Convertible:</b> <math>c</math> can be converted to monotonic or anti-monotonic if items can be properly ordered in processing</li><li><b>Succinct:</b> If the constraint <math>c</math> can be enforced by directly manipulating the data</li></ul>	<ul style="list-style-type: none"><li><b>Data succinct:</b> Data space can be pruned at the initial pattern mining process</li><li><b>Data anti-monotonic:</b> If a transaction <math>t</math> does not satisfy <math>c</math>, then <math>t</math> can be pruned to reduce data processing effort</li></ul>

# How to Handle Multiple Constraints?

---

- It is beneficial to use multiple constraints in pattern mining
- But different constraints may require potentially conflicting item-ordering
  - If there exists conflict ordering between  $c_1$  and  $c_2$ 
    - Try to sort data and enforce *one constraint* first (which one?)
    - Then enforce the other constraint when mining the projected databases
- E.g.  $c_1$ :  $\text{avg}(S.\text{profit}) > 20$ , and  $c_2$ :  $\text{avg}(S.\text{price}) < 50$ 
  - Assume  $c_1$  has more pruning power
  - Sort in profit descending order and use  $c_1$  first
  - For each project DB, sort trans. in price ascending order and use  $c_2$  at mining

# **Chapter 7 : Advanced Frequent Pattern Mining**

---

- ❑ Mining Diverse Patterns
- ❑ Constraint-Based Frequent Pattern Mining
- ❑ Sequential Pattern Mining 
- ❑ Graph Pattern Mining
- ❑ Pattern Mining Application: Mining Software Copy-and-Paste Bugs
- ❑ Summary

# Sequential Pattern Mining

---

- Sequential Pattern and Sequential Pattern Mining
- GSP: Apriori-Based Sequential Pattern Mining
- SPADE: Sequential Pattern Mining in Vertical Data Format
- PrefixSpan: Sequential Pattern Mining by Pattern-Growth
- CloSpan: Mining Closed Sequential Patterns
- Constraint-Based Sequential-Pattern Mining

# Sequential Pattern Mining

---

- What kind of patterns are sequential?
  - Sequential – The order really matters. You can not swap two items in a sequence and have the same sequence.
  - Example: The English language is sequential : Subject -> Verb -> Object.
- 
- Other points:
    - For Sequential Pattern Mining, the time which the items occur is **not** considered.
    - Time Series Analysis does take into account the time in which an item occurred.

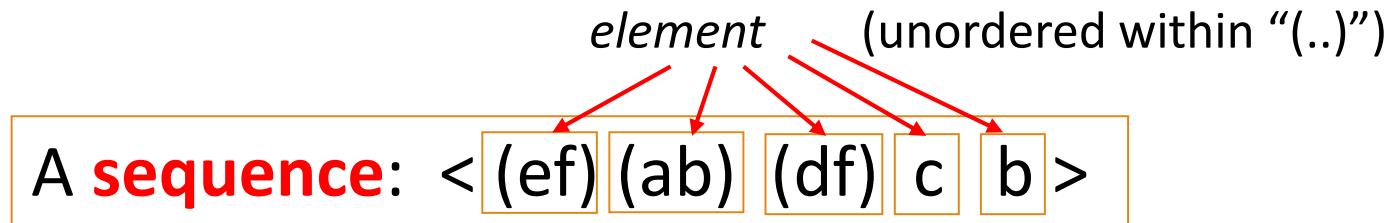
# Sequential Pattern Examples

---

- Application of Sequential pattern Mining
  - **Customer shopping** → Purchase a laptop first, then a digital camera, and then a smartphone.
  - **Medical treatments** → Go to the doctor, get drugs, doctor monitors progress, doctor reacts accordingly -> more/less drugs
  - **Natural disasters** -> Before the disaster, during the disaster, after the disaster.
  - **Scientific Experiments** → Step 1, Step 2, Step 3.
  - **Stocks Markets** → Stocks go up and down together.
  - **Biological sequences, DNA /Protein**→ If you change the order of proteins, it is a different gene.

# Sequential Pattern and Sequential Pattern Mining

- ❑ **Sequential pattern mining:** Given a set of sequences, find the **complete set of frequent subsequences** (i.e., satisfying the min\_sup threshold)



- ❑ An element may contain a set of items (also called events)

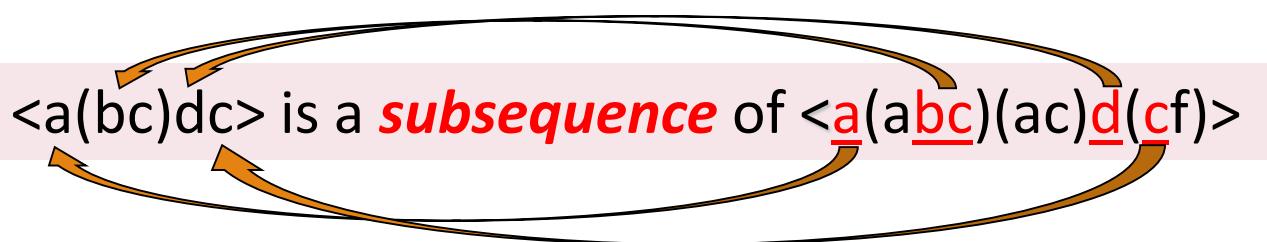
\* Items within an element are **unordered** and we list them alphabetically

A **sequence database**

SID	Sequence
10	<a( <u>abc</u> )(a <u>c</u> )d(cf)>
20	<(ad)c(bc)(ae)>
30	<(ef)( <u>ab</u> )(df) <u>c</u> b>
40	<eg(af)cbc>

# Sequential Pattern and Sequential Pattern Mining

- ❑ **Sequential pattern mining:** Given a set of sequences, find the **complete set of frequent subsequences** (i.e., satisfying the `min_sup` threshold)



A **sequence database**

SID	Sequence
10	$\langle a(\underline{abc})(ac)d(cf) \rangle$
20	$\langle (ad)c(bc)(ae) \rangle$
30	$\langle (ef)(\underline{ab})(df)\underline{cb} \rangle$
40	$\langle eg(af)cbc \rangle$

- ❑ Given support threshold  $min\_sup = 2$ ,  $\langle (\textcolor{red}{ab})c \rangle$  is a sequential pattern

# Sequential Pattern Mining Algorithms

---

- Algorithm requirement: Efficient, scalable, finding complete set, incorporating various kinds of user-specific constraints
- The Apriori property still holds: If a subsequence  $s_1$  is infrequent, none of  $s_1$ 's super-sequences can be frequent
- Representative algorithms
  - GSP (Generalized Sequential Patterns): Srikant & Agrawal @ EDBT'96)
  - Vertical format-based mining: SPADE (Zaki@Machine Learning'00)
  - Pattern-growth methods: PrefixSpan (Pei, et al. @TKDE'04)
- Mining closed sequential patterns: CloSpan (Yan, et al. @SDM'03)
- Constraint-based sequential pattern mining (to be covered in the constraint mining section)

# GSP: Apriori-Based Sequential Pattern Mining

- Initial candidates: All 8-singleton sequences
  - <a>, <b>, <c>, <d>, <e>, <f>, <g>, <h>
- Scan DB once, count support for each candidate

SID	Sequence
10	<(bd)cb(ac)>
20	<(bf)(ce)b(fg)>
30	<(ah)(bf)abf>
40	<(be)(ce)d>
50	<a(bd)bcb(ade)>

$min\_sup = 2$



Cand.	sup
<a>	3
<b>	5
<c>	4
<d>	3
<e>	3
<f>	2
<g>	1
<h>	1

# GSP: Apriori-Based Sequential Pattern Mining

- Example: Generate length-2 candidate sequences

singleton \* singleton – Total:  $(6 * 6)$

*min\_sup = 2*

Cand.	sup
<a>	3
<b>	5
<c>	4
<d>	3
<e>	3
<f>	2
<g>	1
<h>	1



	<a>	<b>	<c>	<d>	<e>	<f>
<a>	<aa>	<ab>	<ac>	<ad>	<ae>	<af>
<b>	<ba>	<bb>	<bc>	<bd>	<be>	<bf>
<c>	<ca>	<cb>	<cc>	<cd>	<ce>	<cf>
<d>	<da>	<db>	<dc>	<dd>	<de>	<df>
<e>	<ea>	<eb>	<ec>	<ed>	<ee>	<ef>
<f>	<fa>	<fb>	<fc>	<fd>	<fe>	<ff>

Sets (unordered) – Total:  $(6 * 5) / 2$

	<a>	<b>	<c>	<d>	<e>	<f>
<a>		<(ab)>	<(ac)>	<(ad)>	<(ae)>	<(af)>
<b>			<(bc)>	<(bd)>	<(be)>	<(bf)>
<c>				<(cd)>	<(ce)>	<(cf)>
<d>					<(de)>	<(df)>
<e>						<(ef)>
<f>						

Apriori Pruning

- w/o pruning  
*(includes g and h):*

$$8 * 8 + 8 * 7 / 2 = 92$$

length-2 candidates

- w/ pruning:  
 $6 * 6 + 6 * 5 / 2 = 51$   
length-2 candidates

# GSP Mining and Pruning

**5<sup>th</sup> scan:** 1 cand. 1 length-5 seq. pat.

<(bd)cba>

length

5

**4<sup>th</sup> scan:** 8 cand. 7 length-4 seq. pat.

<abba> <(bd)bc> ...

4

**3<sup>rd</sup> scan:** 46 cand. 20 length-3 seq. pat. 20  
cand. not in DB at all

<abb> <aab> <aba> **<baa>** <bab> ...

3

**2<sup>nd</sup> scan:** **51** cand. 19 length-2 seq. pat.  
10 cand. not in DB at all

<aa> <ab> ... <af> <ba> <bb> ... <ff> **<(ab)>** ... **<(ef)>**

2

**1<sup>st</sup> scan:** 8 cand. 6 length-1 seq. pat.

<a> <b> <c> <d> <e> <f> **<g>** **<h>**

1

$$6*6 + 6*5/2 = 51$$

- Remove
  - Candidates not in DB**
  - Candidates < min\_sup**

*min\_sup = 2*

SID	Sequence
10	<(bd)cb(ac)>
20	<(bf)(ce)b(fg)>
30	<(ah)(bf)abf>
40	<(be)(ce)d>
50	<a(bd)bcb(ade)>

# GSP Mining and Pruning

---

- Repeat, starting at  $k=1$  until  $k \leq \text{length}$ 
  - Scan DB to find “ $\text{length}-k$ ” frequent sequences
  - Generate “ $\text{length}-(k+1)$ ” candidate sequences from “ $\text{length}-k$ ” frequent sequences using **Apriori**
  - set  $k = k+1$
- Until no frequent sequence or no candidate can be found

**GSP** (Generalized Sequential Patterns): Srikant & Agrawal @ EDBT'96)  
-NOTE: Same team which developed Apriori

# Sequential Pattern Mining in Vertical Data Format: The SPADE Algorithm

- A sequence database is mapped to: <SID, EID>
- Grow the subsequences (patterns) one item at a time by Apriori candidate generation

SID	Sequence
1	<a( <u>abc</u> )(ac)d(cf)>
2	<(ad)c(bc)(ae)>
3	<(ef)(ab)(df) <u>cb</u> >
4	<eg(af)cbc>
<i>min_sup = 2</i>	

Ref: SPADE (Sequential  
Pattern Discovery  
using Equivalent Class)  
[M. Zaki 2001]

SID	EID	Items
1	1	a
1	2	abc
1	3	ac
1	4	d
1	5	cf
2	1	ad
2	2	c
2	3	bc
2	4	ae
3	1	ef
3	2	ab
3	3	df
3	4	c
3	5	b
4	1	e
4	2	g
4	3	af
4	4	c
4	5	b
4	6	c

a		b		...	
SID	EID	SID	EID	...	
1	1	1	2		
1	2	2	3		
1	3	3	2		
2	1	3	5		
2	4	4	5		
3	2				
4	3				

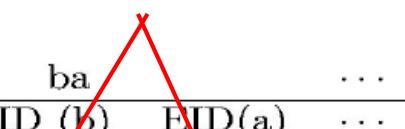
ab		ba		...	
SID	EID (a)	EID(b)	SID	EID (b)	EID(a)
1	1	2	1	2	3
2	1	3	2	3	4
3	2	5			
4	3	5			

aba		...	
SID	EID (a)	EID(b)	EID(a)
1	1	2	3
2	1	3	4

EID (b) < EID (a):  
Corresponds to:  
<a(abc)(ac)d(cf)>

...  
ba  
EID(a)



...  
SID EID (a) EID(b) SID EID (b) EID(a)  
1 1 2 1 2 3  
2 1 3 2 3 4  
3 2 5  
4 3 5

...  
SID EID (a) EID(b) EID(a)  
1 1 2 3  
2 1 3 4

# PrefixSpan: A Pattern-Growth Approach

SID	Sequence
10	<a( <u>abc</u> )(ac)d(cf)>
20	<(ad)c(bc)(ae)>
30	<(ef)( <u>ab</u> )(df) <u>c</u> b>
40	<eg(af)cbc>

min_sup = 2	
Prefix	<u>Suffix (Projection)</u>
<a>	<(abc)(ac)d(cf)>
<aa>	<(_bc)(ac)d(cf)>
<ab>	<(_c)(ac)d(cf)>

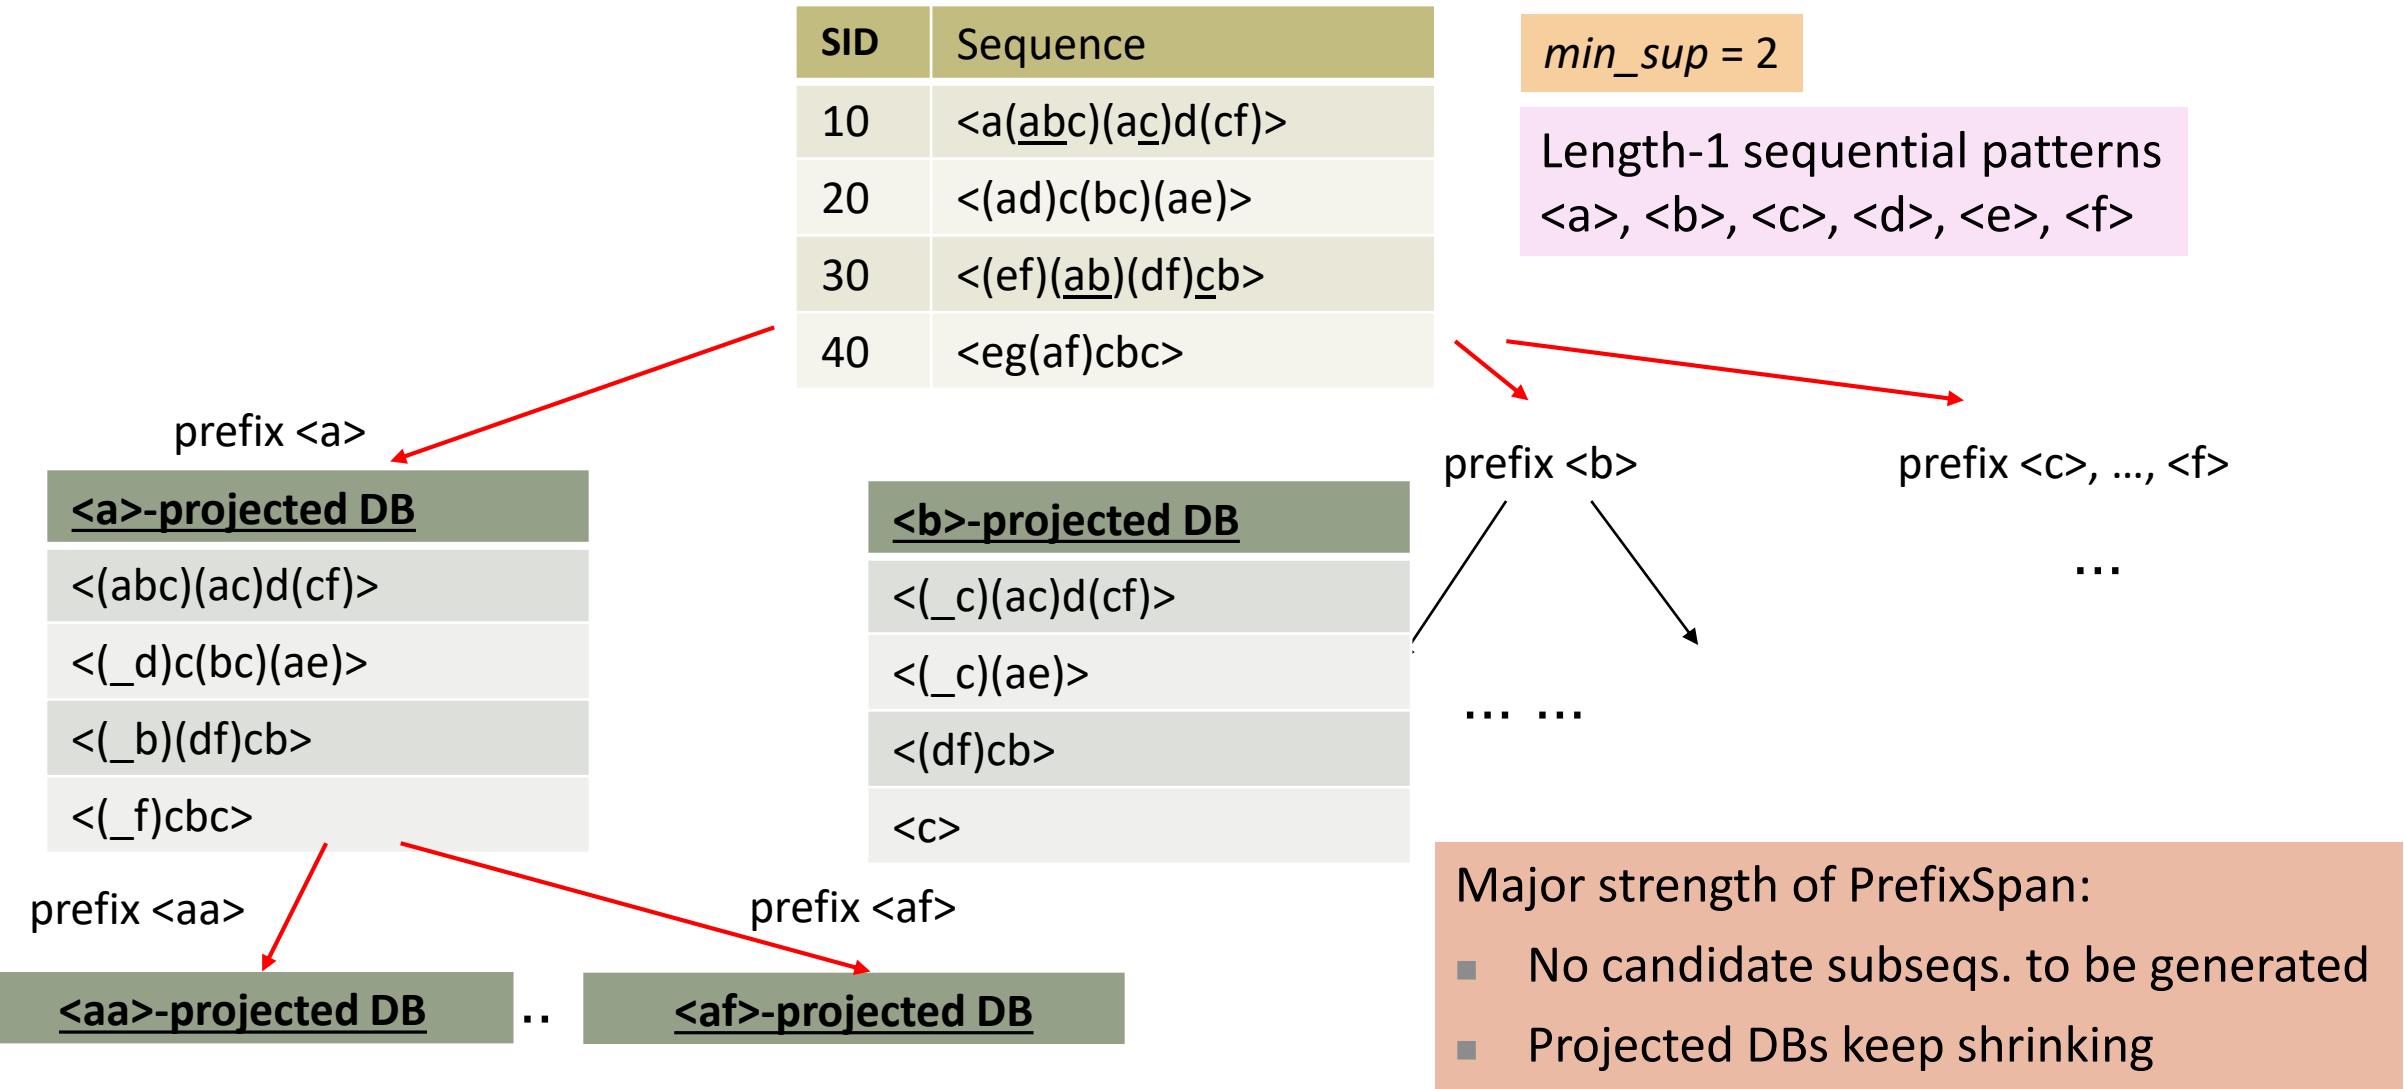
- ❑ PrefixSpan Mining: Prefix Projections
  - ❑ Step 1: Find length-1 sequential patterns
    - ❑ <a>, <b>, <c>, <d>, <e>, <f>
  - ❑ Step 2: Divide search space and mine each projected DB
    - ❑ <a>-projected DB,
    - ❑ <b>-projected DB,
    - ❑ ...
    - ❑ <f>-projected DB, ...

- ❑ Prefix and suffix
  - ❑ Given <a(abc)(ac)d(cf)>
  - ❑ Prefixes: <a>, <aa>, <a(ab)>, <a(abc)>, ...
  - ❑ Suffix: Prefixes-based projection

"\_" is placeholder for prefix

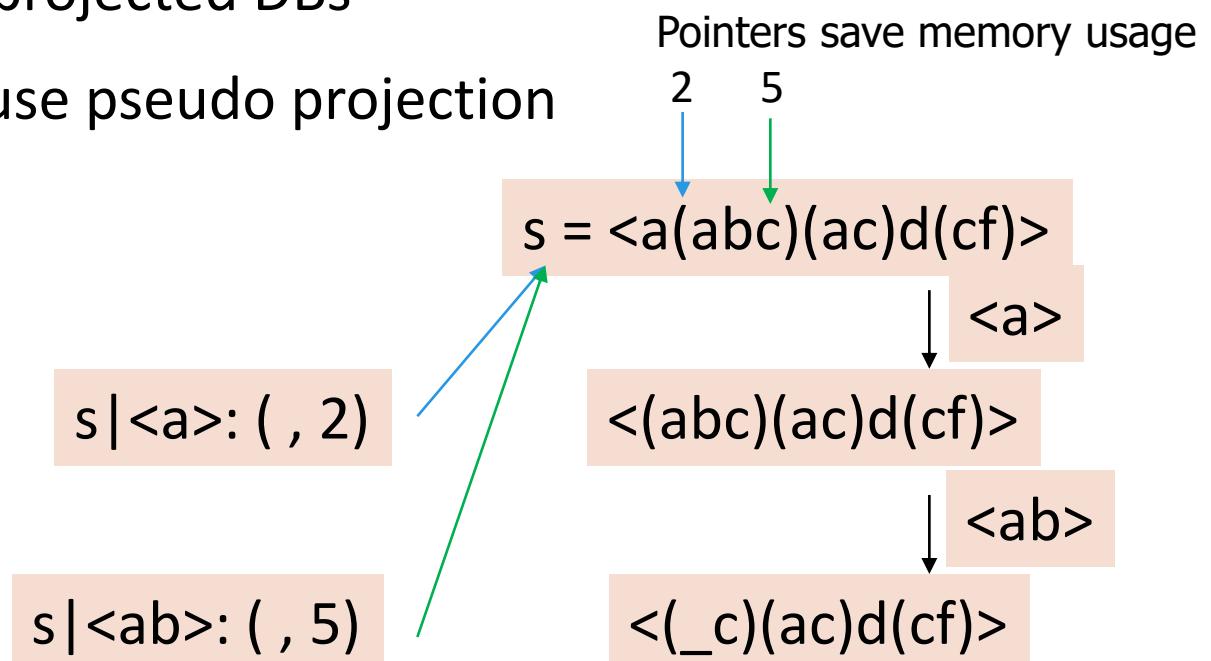
PrefixSpan (Prefix-projected Sequential pattern mining)  
Pei, et al. @TKDE'04

# PrefixSpan: Mining Prefix-Projected DBs



# Implementation Consideration: Pseudo-Projection vs. Physical Projection

- Major cost of PrefixSpan: Constructing projected DBs
  - Suffixes largely repeating in recursive projected DBs
- When DB can be held in main memory, use pseudo projection
  - No physically copying suffixes
  - **Pointer to the sequence**
  - **Offset of the suffix**
- But if it does not fit in memory
  - Physical projection
- Suggested approach:
  - Integration of physical and pseudo-projection
  - Swapping to pseudo-projection when the data fits in memory



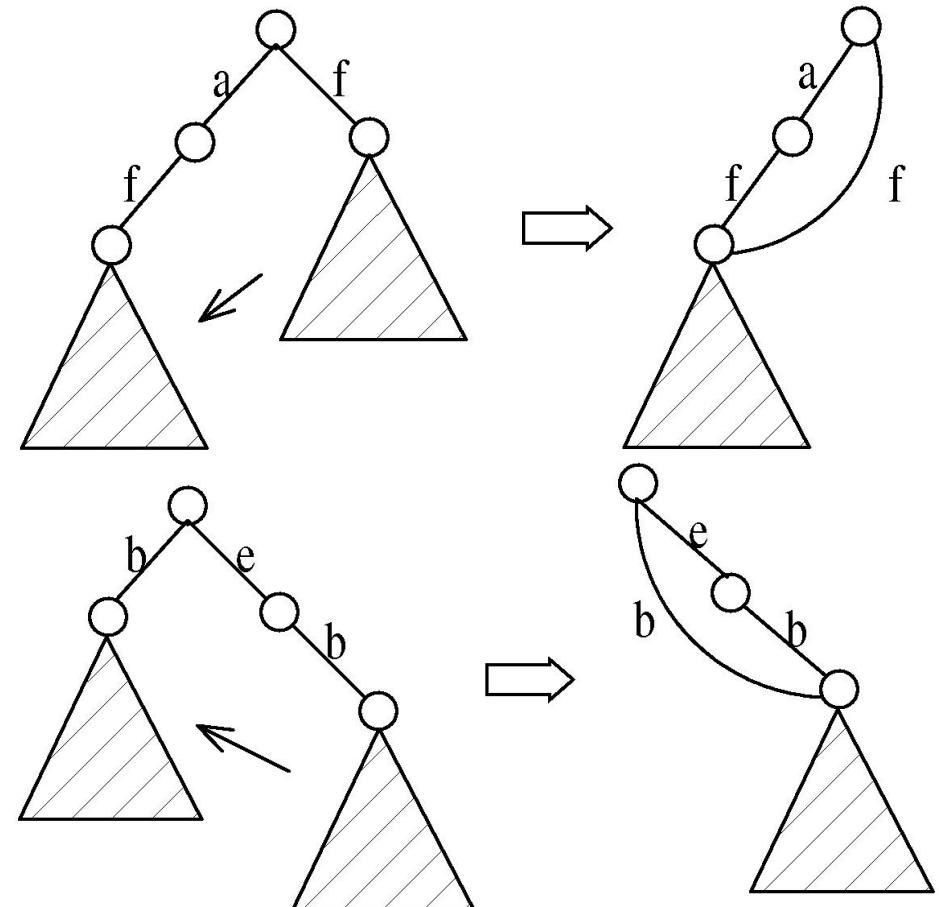
# CloSpan: Mining Closed Sequential Patterns

---

- A **closed sequential pattern**  $s$ : There exists no superpattern  $s'$  such that  $s' \supset s$ , and  $s'$  and  $s$  have the same support
- Which ones are closed?  $\langle abc \rangle: 20, \langle abcd \rangle: 20, \langle abcde \rangle: 15$

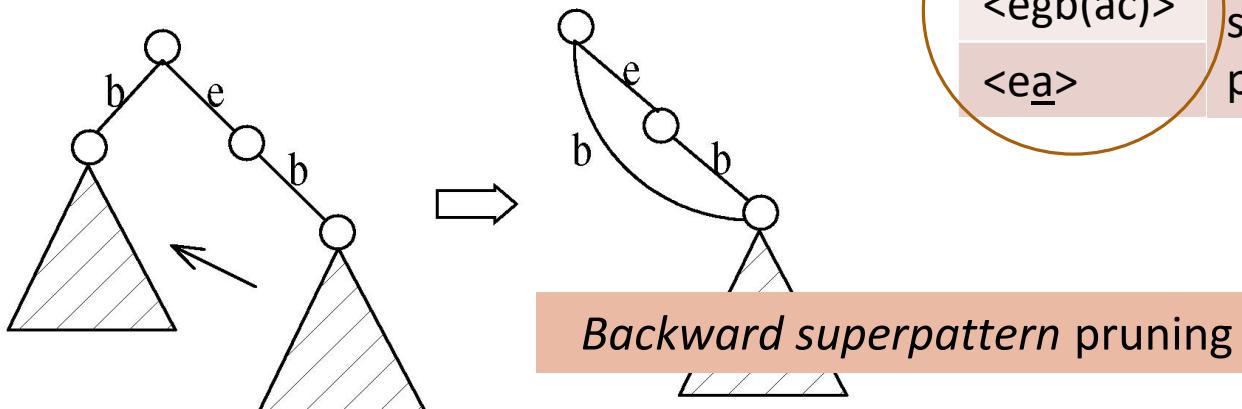
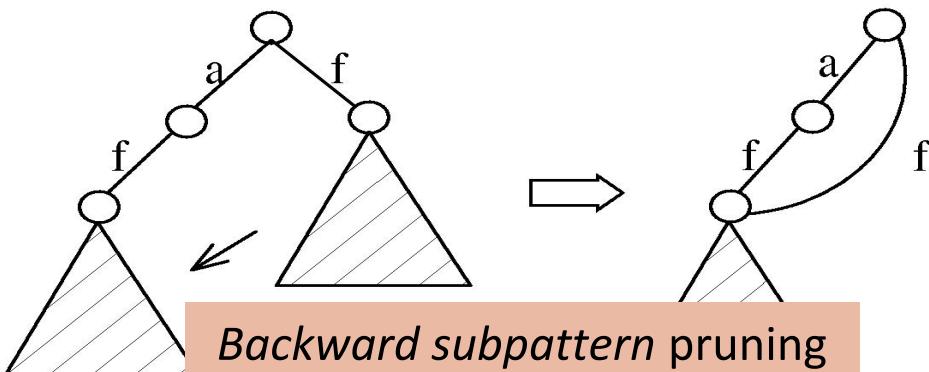
# CloSpan: Mining Closed Sequential Patterns

- Why directly mine closed sequential patterns?
  - Reduce # of (redundant) patterns
  - Attain the same expressive power
- Property P<sub>1</sub>: If  $s \supset s_1$ , s is closed iff two project DBs have the same size
- Explore *Backward Subpattern* and *Backward Superpattern* pruning to prune redundant search space
- Greatly enhances efficiency (Yan, et al., SDM'03)



# CloSpan: When Two Projected DBs Have the Same Size

- If  $s \supset s_1$ ,  $s$  is closed iff two project DBs have the same size
- When two projected sequence DBs have the same size?
- Here is one example:



ID	Sequence	$min\_sup = 2$
1	<aefbcg>	
2	<afegb(ac)>	
3	<(af)ea>	

$\langle a \rangle$      $\langle b \rangle$      $\langle e \rangle$      $\langle f \rangle$

$\langle af \rangle$

$\langle efbcg \rangle$   
 $\langle fegb(ac) \rangle$   
 $\langle (_f)ea \rangle$

$\langle cg \rangle$   
 $\langle ac \rangle$   
size = 6)

$\langle fbcg \rangle$   
 $\langle gb(ac) \rangle$   
 $\langle a \rangle$

$\langle bcg \rangle$   
 $\langle egb(ac) \rangle$   
 $\langle ea \rangle$

Only need to keep  
size = 12 (including  
parentheses)

$\langle cg \rangle$   
 $\langle ac \rangle$

$\langle b \rangle$

$\langle b \rangle$

$\langle bcg \rangle$   
 $\langle egb(ac) \rangle$   
 $\langle ea \rangle$

# Constraint-Based Sequential-Pattern Mining

---

- Share many similarities with constraint-based itemset mining
- **Anti-monotonic:** If  $S$  violates  $c$ , the super-sequences of  $S$  also violate  $c$ 
  - $\text{sum}(S.\text{price}) < 150; \text{min}(S.\text{value}) > 10$
- **Monotonic:** If  $S$  satisfies  $c$ , the super-sequences of  $S$  also do so
  - $\text{element\_count}(S) > 5; S \supseteq \{\text{PC}, \text{digital\_camera}\}$
- **Data anti-monotonic:** If a sequence  $s_1$  with respect to  $S$  violates  $c_3$ ,  $s_1$  can be removed
  - $c_3: \text{sum}(S.\text{price}) \geq v$
- **Succinct:** Enforce constraint  $c$  by explicitly manipulating data
  - $S \supseteq \{\text{i-phone}, \text{MacAir}\}$
- **Convertible:** Projection based on the sorted value not sequence order
  - $\text{value\_avg}(S) < 25; \text{profit\_sum}(S) > 160$
  - $\text{max}(S)/\text{avg}(S) < 2; \text{median}(S) - \text{min}(S) > 5$

# Timing-Based Constraints in Seq.-Pattern Mining

---

- **Order constraint:** Some items must happen before the other
  - {algebra, geometry} → {calculus} (where “→” indicates ordering)
  - Anti-monotonic: Constraint-violating sub-patterns pruned
- **Min-gap/max-gap constraint:** Confines two elements in a pattern
  - E.g., mingap = 1, maxgap = 4
  - Succinct: Enforced directly during pattern growth
- **Max-span constraint:** Maximum allowed time difference between the 1<sup>st</sup> and the last elements in the pattern
  - E.g., maxspan (S) = 60 (days)
  - Succinct: Enforced directly when the 1<sup>st</sup> element is determined
- **Window size constraint:** Events in an element do not have to occur at the same time: Enforce max allowed time difference
  - E.g., window-size = 2: Various ways to merge events into elements

# Episodes and Episode Pattern Mining

- ❑ Episodes and regular expressions: Alternative to seq. patterns
  - ❑ Serial episodes: AB       a total order relationship: first A then B
  - ❑ Parallel episodes: A|B       a partial order relationship: A and B can be in any order
  - ❑ Regular expressions: (A|B)C\*(DE)       (DE) means D, E happen in the same time window
- ❑ E.g. Given a large shopping sequence database, one may like to find
  - ❑ Suppose the pattern order follows the template (A|B)C\*(D E), and
  - ❑ Sum of the prices of A, B, C\*, D, and E is greater than \$100, where C\* means C appears \*-times
  - ❑ How to efficiently mine such episode patterns?

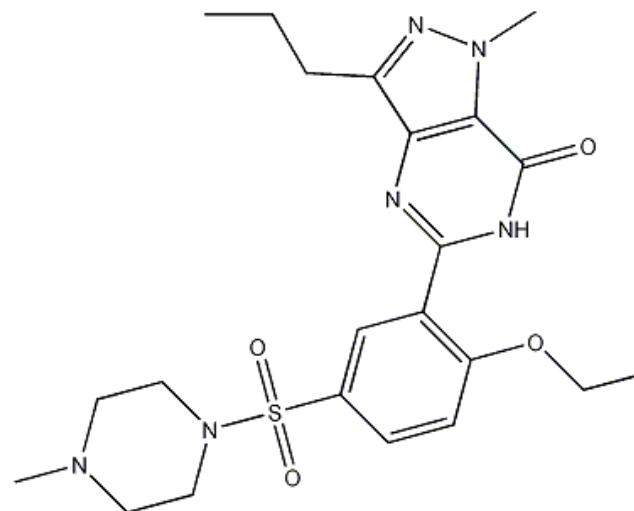
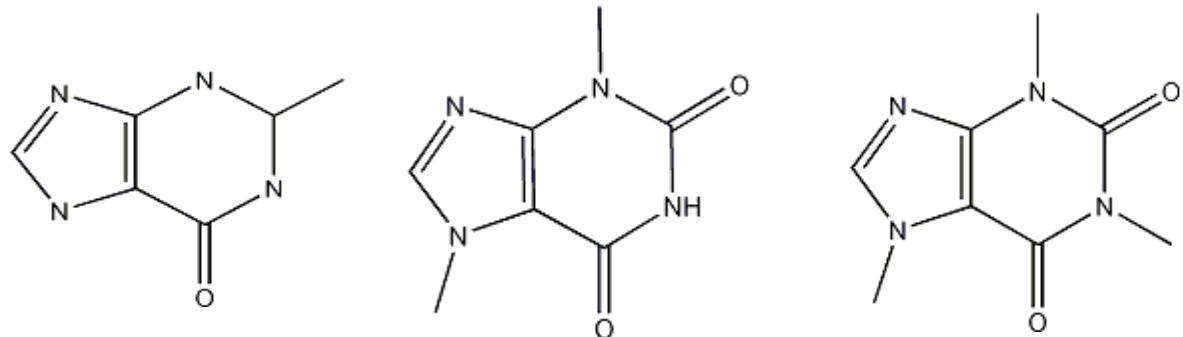
# **Chapter 7 : Advanced Frequent Pattern Mining**

---

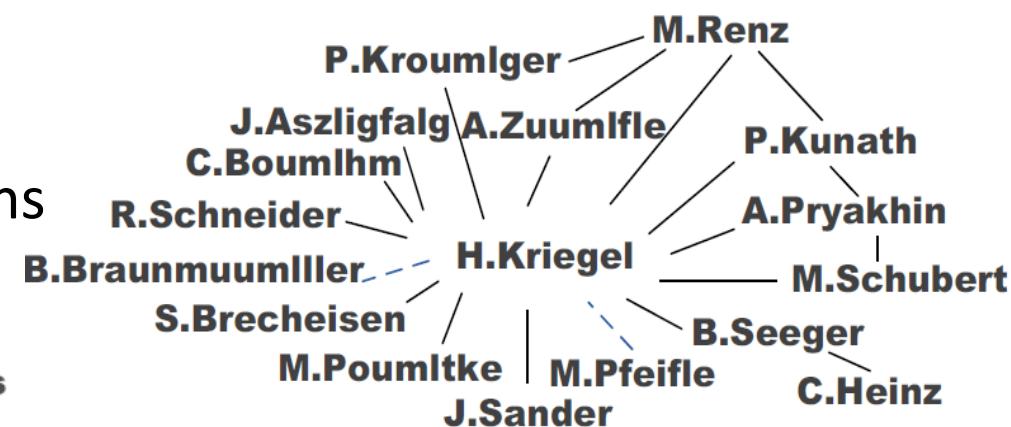
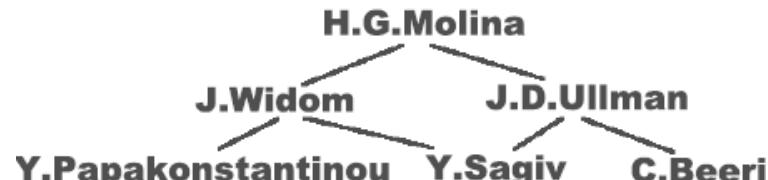
- Mining Diverse Patterns
- Constraint-Based Frequent Pattern Mining
- Sequential Pattern Mining
- Graph Pattern Mining 
- Pattern Mining Application: Mining Software Copy-and-Paste Bugs
- Summary

# What Is Graph Pattern Mining?

- Chem-informatics:
  - Mining frequent chemical compound structures

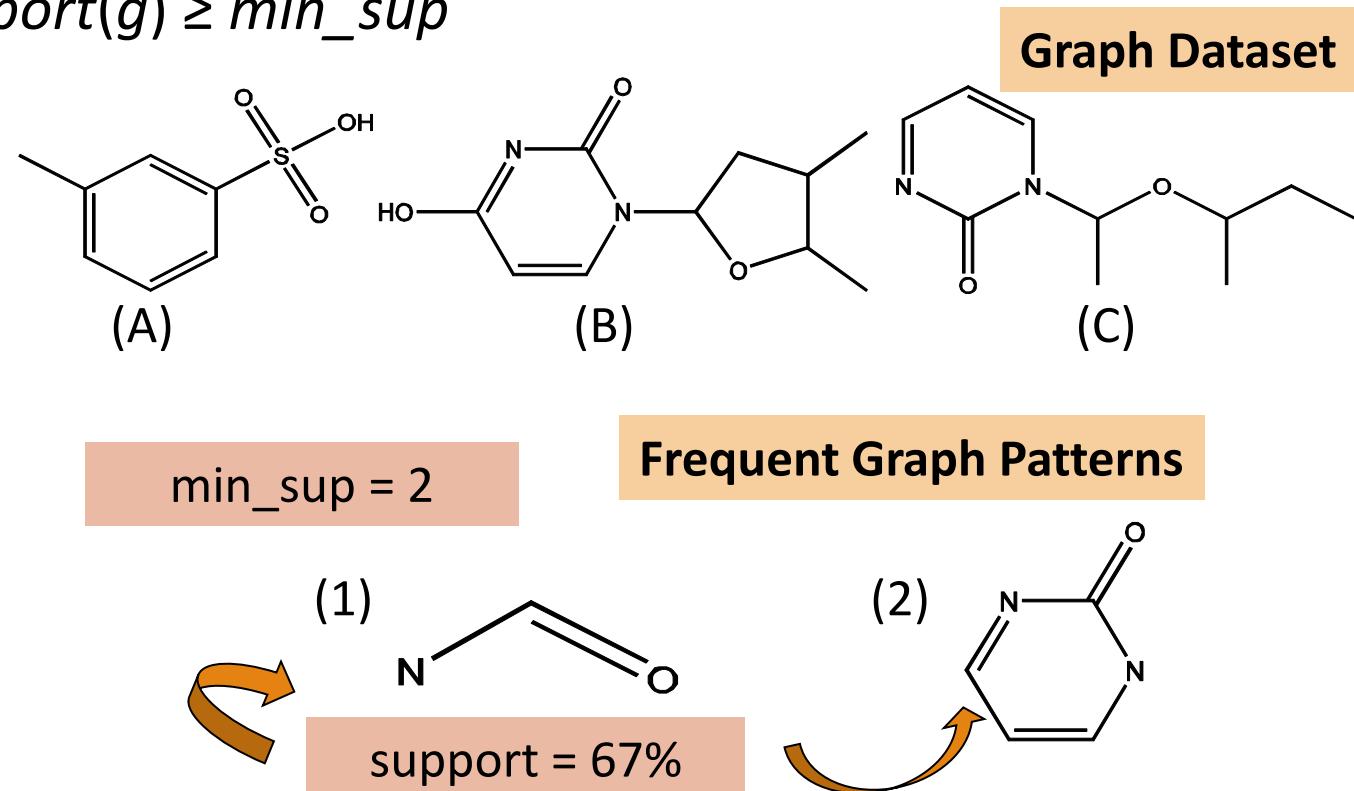


- Social networks, web communities, tweets, ...
  - Finding frequent research collaboration subgraphs



# Frequent (Sub)Graph Patterns

- Given a labeled graph dataset  $D = \{G_1, G_2, \dots, G_n\}$ , the supporting graph set of a subgraph  $g$  is  $D_g = \{G_i \mid g \subseteq G_i, G_i \in D\}$ 
  - $\text{support}(g) = |D_g| / |D|$
- A (sub)graph  $g$  is *frequent* if  $\text{support}(g) \geq \text{min\_sup}$
- Ex.: Chemical structures
- Alternative:
  - Mining frequent subgraph patterns from a single large graph or network



# **Applications of Graph Pattern Mining**

---

- Bioinformatics
  - Gene networks, protein interactions, metabolic pathways
- Chem-informatics: Mining chemical compound structures
- Social networks, web communities, tweets, ...
- Cell phone networks, computer networks, ...
- Web graphs, XML structures, Semantic Web, information networks
- Software engineering: Program execution flow analysis
- Building blocks for graph classification, clustering, compression, comparison, and correlation analysis
- Graph indexing and graph similarity search

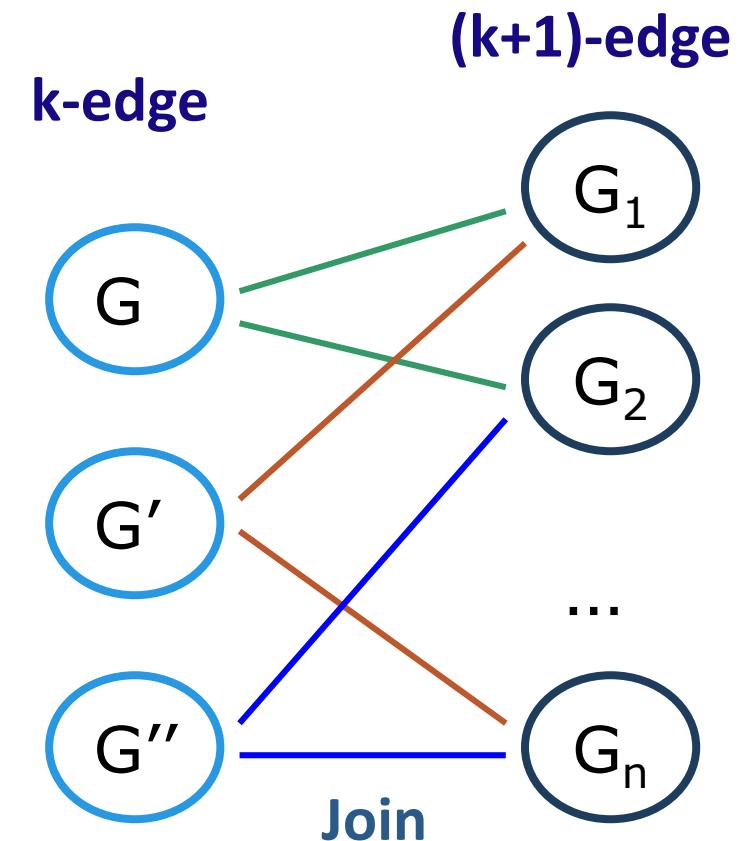
# Graph Pattern Mining Algorithms: Different Methodologies

---

- Generation of candidate subgraphs
  - Apriori vs. pattern growth (e.g., FSG vs. gSpan)
- Search order
  - Breadth vs. depth
- Elimination of duplicate subgraphs
  - Passive vs. active (e.g., gSpan [Yan & Han, 2002])
- Support calculation
  - Store embeddings (e.g., GASTON [Nijssen & Kok, 2004], FFSM [Huan, Wang, & Prins, 2003], MoFa [Borgelt & Berthold, ICDM'02])
- Order of pattern discovery
  - Path → tree → graph (e.g., GASTON [Nijssen & Kok, 2004])

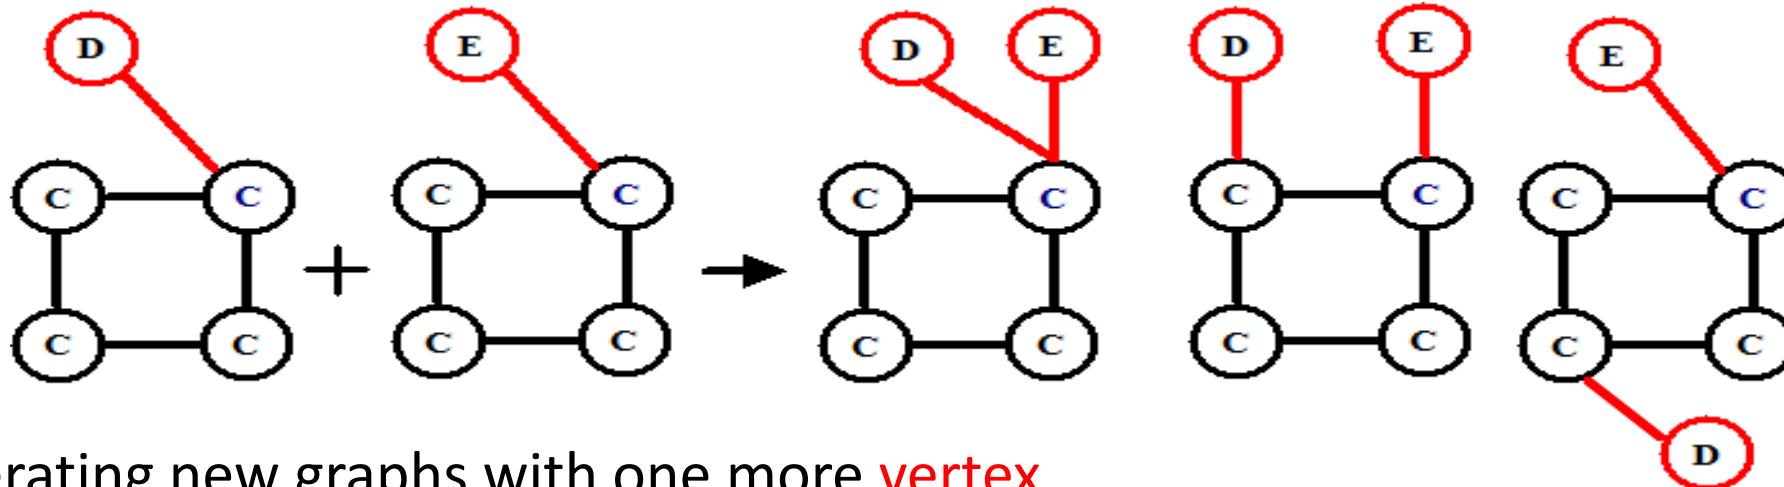
# Apriori-Based Approach

- The Apriori property (anti-monotonicity): A size- $k$  subgraph is **frequent** if and only if all of its **subgraphs are frequent**
- A candidate size- $(k+1)$  edge/vertex subgraph is generated if its corresponding two  $k$ -edge/vertex subgraphs are frequent
- Iterative mining process:
  - Candidate-generation → candidate pruning → support counting → candidate elimination



# Candidate Generation: Vertex Growing vs. Edge Growing

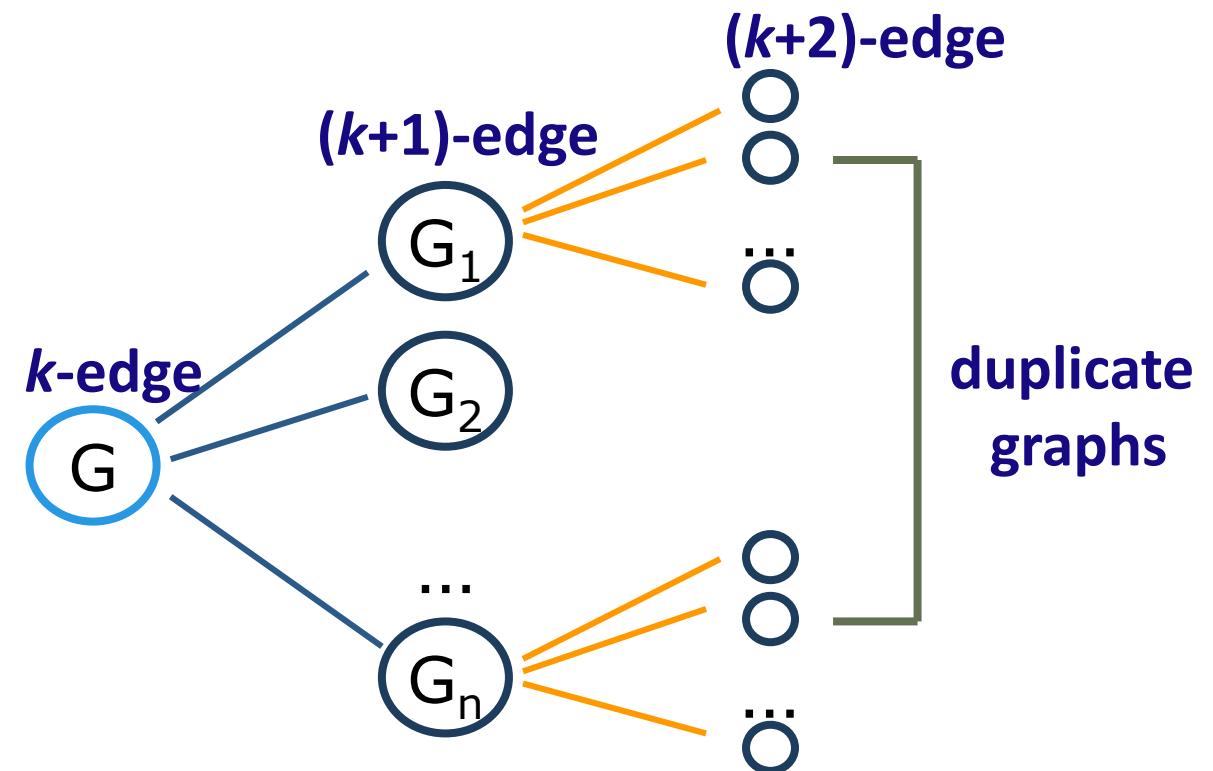
- Methodology: Breadth-search, Apriori joining two size- $k$  graphs
  - Many possibilities at generating size- $(k+1)$  candidate graphs



- Generating new graphs with one more **vertex**
  - AGM (Inokuchi, Washio, & Motoda, PKDD'00)
- Generating new graphs with one more **edge**
  - FSG (Kuramochi & Karypis, ICDM'01)
- Performance shows *via edge growing* is more efficient

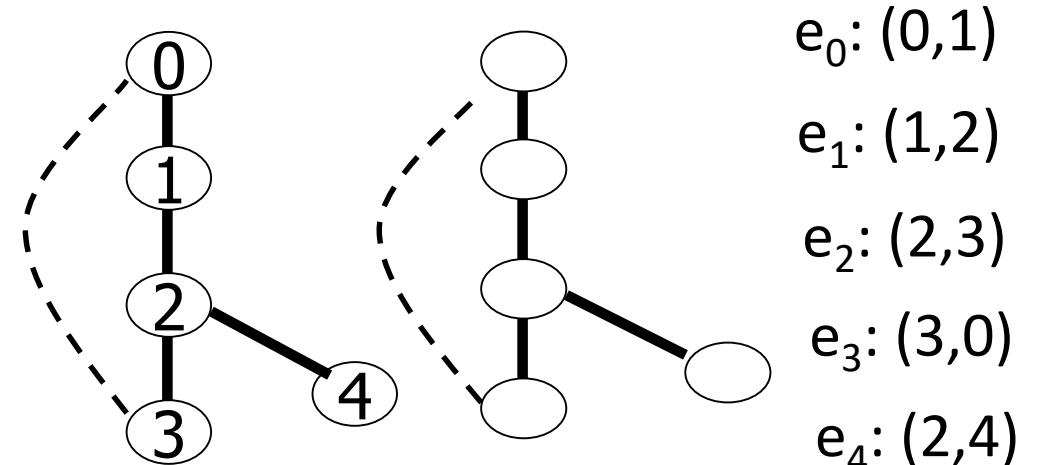
# Pattern-Growth Approach

- Depth-first growth of subgraphs from  $k$ -edge to  $(k+1)$ -edge, then  $(k+2)$ -edge subgraphs
- Major challenge
  - Generating many duplicate subgraphs
- Major idea to solve the problem
  - Define an order to generate subgraphs
  - DFS spanning tree: Flatten a graph into a sequence using depth-first search
  - gSpan (Yan & Han, ICDM'02)



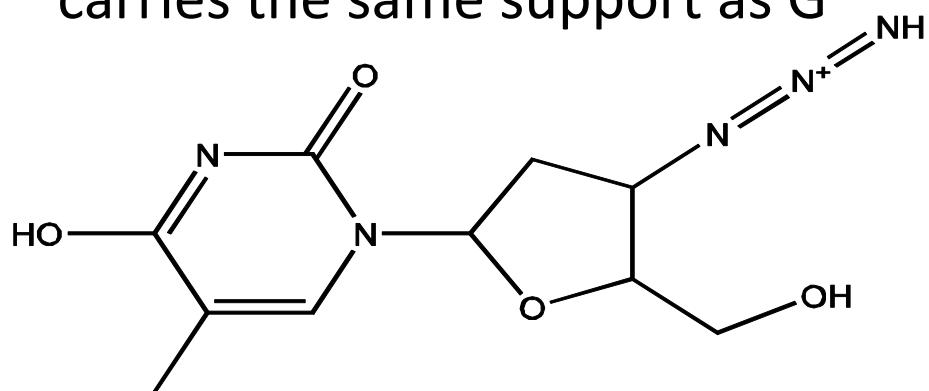
# gSPAN: Graph Pattern Growth in Order

- **Right-most path extension** in subgraph pattern growth
- Right-most path: The path from root to the right-most leaf (choose the vertex with the **smallest** index at each step)
- Reduce generation of duplicate subgraphs
- **Completeness:** The enumeration of graphs using right-most path extension is complete
- DFS code: Flatten a graph into a sequence using depth-first search



# Why Mine Closed Graph Patterns?

- $2^n$  subgraphs  $\rightarrow$  *closed frequent subgraphs*
- A frequent graph  $G$  is *closed* if there exists no supergraph of  $G$  that carries the same support as  $G$

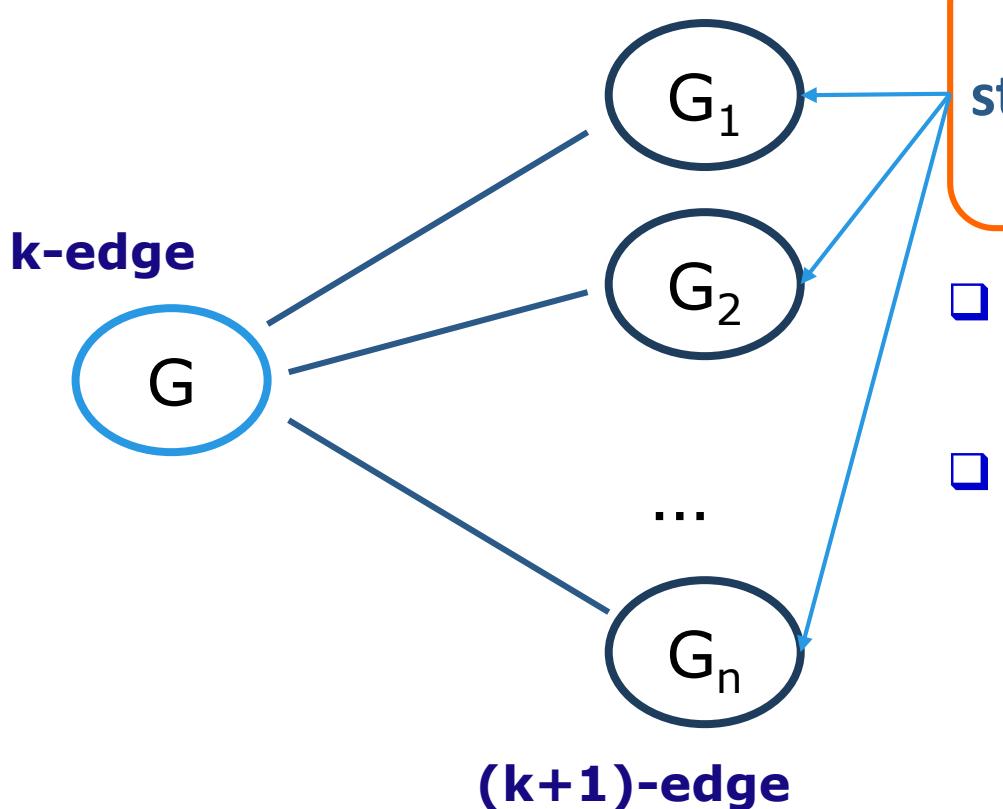


If this subgraph is *closed* in the graph dataset, it implies that none of its frequent super-graphs carries the same support

- *Lossless compression*: Does not contain non-closed graphs, but still ensures that the mining result is complete
- Algorithm CloseGraph: Mines closed graph patterns directly

# CloseGraph: Directly Mining Closed Graph Patterns

- CloseGraph: Mining closed graph patterns by extending gSpan (Yan & Han, KDD'03)

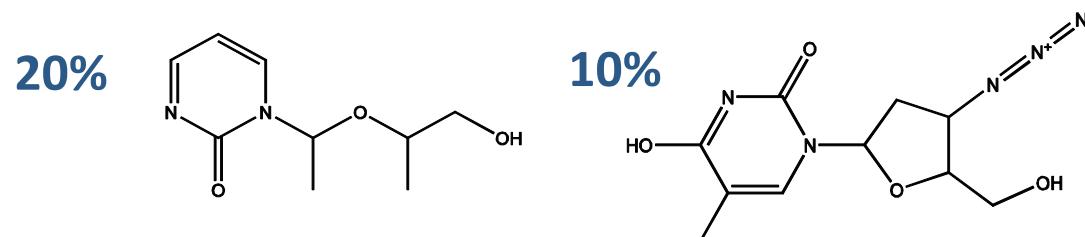


At what condition can we stop searching their children, i.e., early termination?

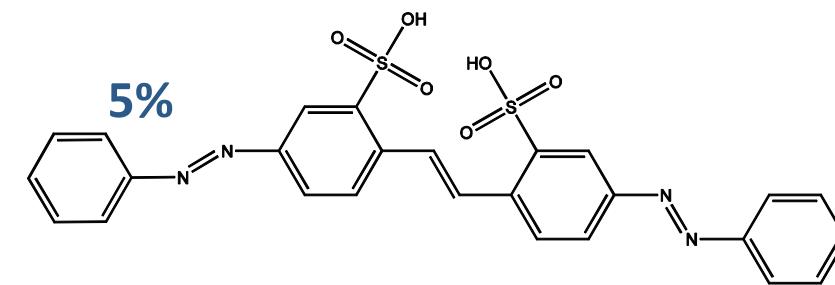
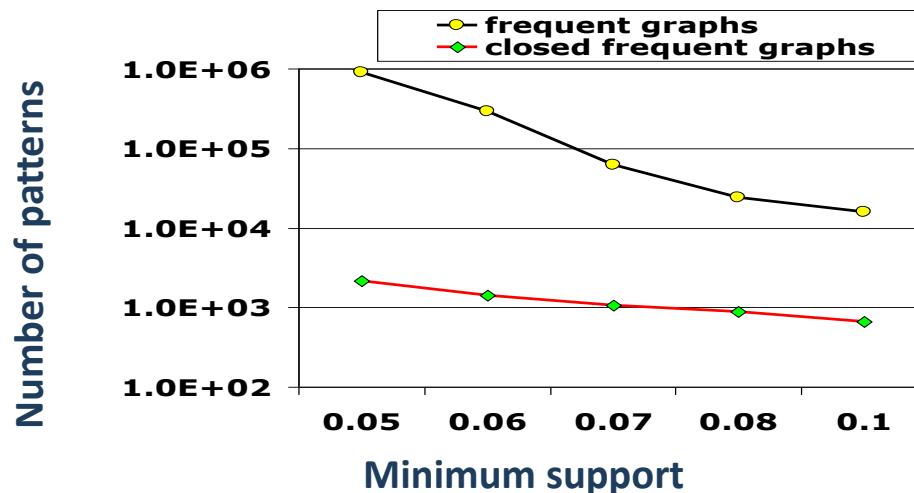
- Suppose  $G$  and  $G_1$  are frequent, and  $G$  is a subgraph of  $G_1$
- If **in any part of the graph in the dataset where  $G$  occurs,  $G_1$  also occurs**, then we need not grow  $G$  (except some special, subtle cases), since none of  $G$ 's children will be closed except those of  $G_1$

# Experiment and Performance Comparison

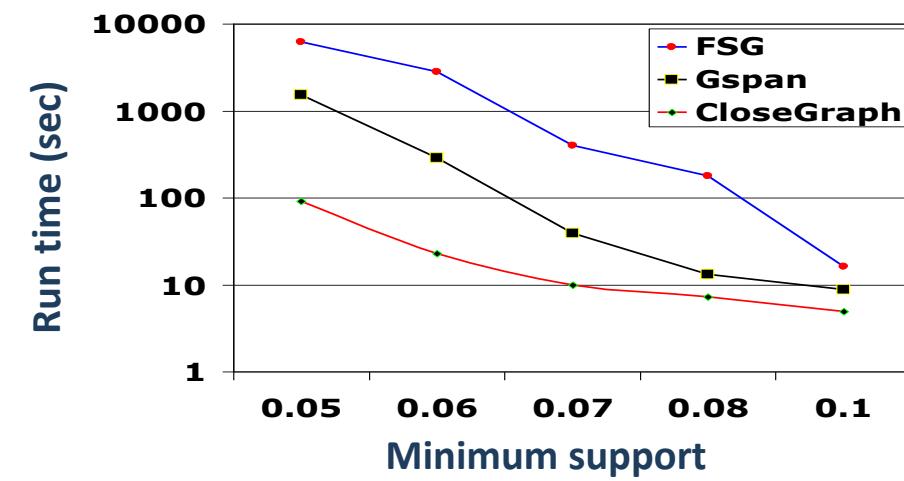
- The AIDS antiviral screen compound dataset from NCI/NIH
- The dataset contains 43,905 chemical compounds
- Discovered patterns: The smaller minimum support, the bigger and more interesting subgraph patterns discovered



# of Patterns: Frequent vs. Closed



Runtime: Frequent vs. Closed



# **Chapter 7 : Advanced Frequent Pattern Mining**

---

- ❑ Mining Diverse Patterns
- ❑ Constraint-Based Frequent Pattern Mining
- ❑ Sequential Pattern Mining
- ❑ Graph Pattern Mining
- ❑ Pattern Mining Application: Mining Software Copy-and-Paste Bugs 
- ❑ Summary

# Pattern Mining Application: Software Bug Detection

---

- **Mining rules from source code**
  - Bugs as deviant behavior (e.g., by statistical analysis)
  - Mining programming rules (e.g., by frequent itemset mining)
  - Mining function precedence protocols (e.g., by frequent subsequence mining)
  - Revealing neglected conditions (e.g., by frequent itemset/subgraph mining)
- **Mining rules from revision histories**
  - By frequent itemset mining
- **Mining copy-paste patterns from source code**
  - Find copy-paste bugs (e.g., CP-Miner [Li et al., OSDI'04]) (to be discussed here)
    - Reference: Z. Li, S. Lu, S. Myagmar, Y. Zhou, “[CP-Miner](#): A Tool for Finding Copy-paste and Related Bugs in Operating System Code”, OSDI’04

# Application Example: Mining Copy-and-Paste Bugs

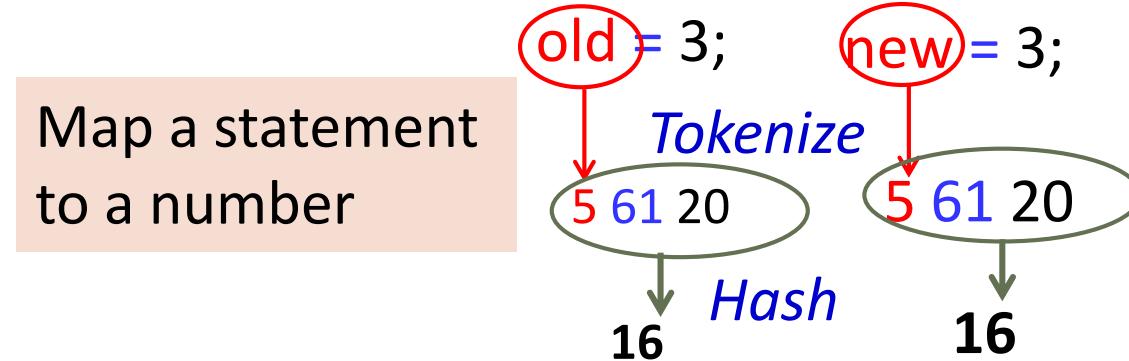
- ❑ Copy-pasting is common
  - ❑ 12% in Linux file system
  - ❑ 19% in X Window system
- ❑ Copy-pasted code is error-prone
- ❑ Mine “*forget-to-change*” bugs by sequential pattern mining
  - ❑ Build a sequence database from source code
  - ❑ Mining sequential patterns
  - ❑ Finding mismatched identifier names & bugs

```
void __init prom_meminit(void)
{
    .....
    for (i=0; i<n; i++) {
        total[i].adr = list[i].addr;
        total[i].bytes = list[i].size;
        total[i].more = &total[i+1];
    }
    .....
    for (i=0; i<n; i++) {
        taken[i].adr = list[i].addr;
        taken[i].bytes = list[i].size,
        taken[i].more = &total[i+1];
    }
}
```

Code copy-and-pasted but **forget to change “id”!**

# Building Sequence Database from Source Code

- Statement  $\xrightarrow{\text{(mapped to)}}$  number
- Tokenize each component
  - Different operators, constants, key words  $\rightarrow$  different tokens
  - Same type of identifiers  $\rightarrow$  same token
- Program  $\rightarrow$  A long sequence
- Cut the long sequence by blocks



Hash values

65
16
16
71
...
65
16
16
71

```
for (i=0; i<n; i++) {  
    total[i].adr = list[i].addr;  
    total[i].bytes = list[i].size;  
    total[i].more = &total[i+1];  
}
```

.....

```
for (i=0; i<n; i++) {  
    taken[i].adr = list[i].addr;  
    taken[i].bytes = list[i].size;  
    taken[i].more = &total[i+1];  
}
```

Final sequence DB:

(65)  
(16, 16, 71)  
...  
(65)  
(16, 16, 71)

# Sequential Pattern Mining & Detecting “Forget-to-Change” Bugs

## □ Modification to the *sequence pattern mining algorithm*

- Constrain the max gap

(16, 16, 71)

.....

(16, 16, 10, 71)

Allow a maximal gap:  
inserting statements  
in copy-and-paste

## □ Composing Larger Copy-Pasted Segments

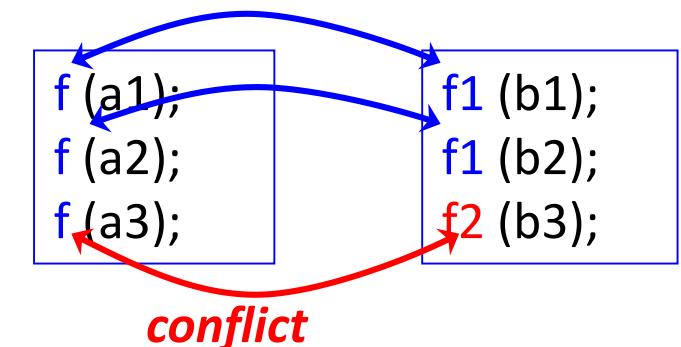
- Combine the neighboring copy-pasted segments repeatedly

## □ Find conflicts: Identify names that cannot be mapped to the corresponding ones

- E.g., 1 out of 4 “**total**” is unchanged, *unchanged ratio* = 0.25

- If  $0 < \text{unchanged ratio} < \text{threshold}$ , then report it as a bug

- CP-Miner reported many C-P bugs in Linux, Apache, ... out of millions of LOC (lines of code)



# **Chapter 7 : Advanced Frequent Pattern Mining**

---

- ❑ Mining Diverse Patterns
- ❑ Constraint-Based Frequent Pattern Mining
- ❑ Sequential Pattern Mining
- ❑ Graph Pattern Mining
- ❑ Pattern Mining Application: Mining Software Copy-and-Paste Bugs
- ❑ Summary 

# Summary: Advanced Frequent Pattern Mining

---

- ❑ Mining Diverse Patterns
  - ❑ Mining Multiple-Level Associations
  - ❑ Mining Multi-Dimensional Associations
  - ❑ Mining Quantitative Associations
  - ❑ Mining Negative Correlations
  - ❑ Mining Compressed and Redundancy-Aware Patterns
- ❑ Sequential Pattern Mining
  - ❑ Sequential Pattern and Sequential Pattern Mining
  - ❑ GSP: Apriori-Based Sequential Pattern Mining
  - ❑ SPADE: Sequential Pattern Mining in Vertical Data Format
  - ❑ PrefixSpan: Sequential Pattern Mining by Pattern-Growth
  - ❑ CloSpan: Mining Closed Sequential Patterns
- ❑ Constraint-Based Frequent Pattern Mining
  - ❑ Why Constraint-Based Mining?
  - ❑ Constrained Mining with Pattern Anti-Monotonicity
  - ❑ Constrained Mining with Pattern Monotonicity
  - ❑ Constrained Mining with Data Anti-Monotonicity
  - ❑ Constrained Mining with Succinct Constraints
  - ❑ Constrained Mining with Convertible Constraints
  - ❑ Handling Multiple Constraints
  - ❑ Constraint-Based Sequential-Pattern Mining
- ❑ Graph Pattern Mining
  - ❑ Graph Pattern and Graph Pattern Mining
  - ❑ Apriori-Based Graph Pattern Mining Methods
  - ❑ gSpan: A Pattern-Growth-Based Method
  - ❑ CloseGraph: Mining Closed Graph Patterns
- ❑ Pattern Mining Application: Mining Software Copy-and-Paste Bugs

# References: Mining Diverse Patterns

---

- R. Srikant and R. Agrawal, "Mining generalized association rules", VLDB'95
- Y. Aumann and Y. Lindell, "A Statistical Theory for Quantitative Association Rules", KDD'99
- K. Wang, Y. He, J. Han, "Pushing Support Constraints Into Association Rules Mining", IEEE Trans. Knowledge and Data Eng. 15(3): 642-658, 2003
- D. Xin, J. Han, X. Yan and H. Cheng, "On Compressing Frequent Patterns", Knowledge and Data Engineering, 60(1): 5-29, 2007
- D. Xin, H. Cheng, X. Yan, and J. Han, "Extracting Redundancy-Aware Top-K Patterns", KDD'06
- J. Han, H. Cheng, D. Xin, and X. Yan, "Frequent Pattern Mining: Current Status and Future Directions", Data Mining and Knowledge Discovery, 15(1): 55-86, 2007
- F. Zhu, X. Yan, J. Han, P. S. Yu, and H. Cheng, "Mining Colossal Frequent Patterns by Core Pattern Fusion", ICDE'07

# References: Constraint-Based Frequent Pattern Mining

---

- R. Srikant, Q. Vu, and R. Agrawal, “Mining association rules with item constraints”, KDD'97
- R. Ng, L.V.S. Lakshmanan, J. Han & A. Pang, “Exploratory mining and pruning optimizations of constrained association rules”, SIGMOD'98
- G. Grahne, L. Lakshmanan, and X. Wang, “Efficient mining of constrained correlated sets”, ICDE'00
- J. Pei, J. Han, and L. V. S. Lakshmanan, “Mining Frequent Itemsets with Convertible Constraints”, ICDE'01
- J. Pei, J. Han, and W. Wang, “Mining Sequential Patterns with Constraints in Large Databases”, CIKM'02
- F. Bonchi, F. Giannotti, A. Mazzanti, and D. Pedreschi, “ExAnte: Anticipated Data Reduction in Constrained Pattern Mining”, PKDD'03
- F. Zhu, X. Yan, J. Han, and P. S. Yu, “gPrune: A Constraint Pushing Framework for Graph Pattern Mining”, PAKDD'07

# References: Sequential Pattern Mining

---

- R. Srikant and R. Agrawal, "Mining sequential patterns: Generalizations and performance improvements", EDBT'96
- M. Zaki, "SPADE: An Efficient Algorithm for Mining Frequent Sequences", Machine Learning, 2001
- J. Pei, J. Han, B. Mortazavi-Asl, J. Wang, H. Pinto, Q. Chen, U. Dayal, and M.-C. Hsu, "Mining Sequential Patterns by Pattern-Growth: The PrefixSpan Approach", IEEE TKDE, 16(10), 2004
- X. Yan, J. Han, and R. Afshar, "CloSpan: Mining Closed Sequential Patterns in Large Datasets", SDM'03
- J. Pei, J. Han, and W. Wang, "Constraint-based sequential pattern mining: the pattern-growth methods", J. Int. Inf. Sys., 28(2), 2007
- M. N. Garofalakis, R. Rastogi, K. Shim: Mining Sequential Patterns with Regular Expression Constraints. IEEE Trans. Knowl. Data Eng. 14(3), 2002
- H. Mannila, H. Toivonen, and A. I. Verkamo, "Discovery of frequent episodes in event sequences", Data Mining and Knowledge Discovery, 1997

# References: Graph Pattern Mining

---

- C. Borgelt and M. R. Berthold, Mining molecular fragments: Finding relevant substructures of molecules, ICDM'02
- J. Huan, W. Wang, and J. Prins. Efficient mining of frequent subgraph in the presence of isomorphism, ICDM'03
- A. Inokuchi, T. Washio, and H. Motoda. An apriori-based algorithm for mining frequent substructures from graph data, PKDD'00
- M. Kuramochi and G. Karypis. Frequent subgraph discovery, ICDM'01
- S. Nijssen and J. Kok. A Quickstart in Frequent Structure Mining can Make a Difference. KDD'04
- N. Vanetik, E. Gudes, and S. E. Shimony. Computing frequent graph patterns from semistructured data, ICDM'02
- X. Yan and J. Han, gSpan: Graph-Based Substructure Pattern Mining, ICDM'02
- X. Yan and J. Han, CloseGraph: Mining Closed Frequent Graph Patterns, KDD'03
- X. Yan, P. S. Yu, J. Han, Graph Indexing: A Frequent Structure-based Approach, SIGMOD'04
- X. Yan, P. S. Yu, and J. Han, Substructure Similarity Search in Graph Databases, SIGMOD'05

